

Maulana Azad National Institute of Technology

(An Institute of National Importance)

Bhopal – 462003 (India)



Department of Computer Science & Engineering

LAB ASSIGNMENT

HADOOP/CUDA (CSE-328)

Name: Jishan Shaikh (161112013)

Instructor: Rahul Singh

Date: March 7, 2019 (Mid Submission)

Index

CUDA:

1. Write a CUDA Program for Parallel Vector addition.
2. Write a CUDA Program for Parallel Matrix Multiplication.
3. Write a CUDA Program for Parallel Binary Search.
4. Write a CUDA Program for finding minimum element in an array using Parallel algorithm.
5. Write a CUDA Program for finding K-shortest path for a given graph with maximum possible parallelism.

HADOOP:

1. Write a Map Reduce program to calculate word frequency in given document.
2. Calculate average marks of every student in the given dataset.
3. Write a program to illustrate use of combiner in Map Reduce.
4. Write a program to illustrate use of partitioner in Map Reduce.
5. Write a Map Reduce Program to return number of unique words in given document.

CUDA:

Write a CUDA Program for Parallel Vector addition.

Program:

```
// Program for Parallel Vector Addition in CUDA
// For Hadoop-CUDA Lab

#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>
#include <time.h>

#define N 1024    // size of array

__global__ void add(int *a,int *b, int *c) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid < N){
        c[tid] = a[tid]+b[tid];
    }
}

int main(int argc, char *argv[]) {
    int T = 10, B = 1;    // threads per block and blocks per grid, taking default values
    int a[N],b[N],c[N];
    int *dev_a, *dev_b, *dev_c;

    printf("Size of array = %d\n", N);
    do {
        printf("Enter number of threads per block: ");
        scanf("%d",&T);
        printf("\nEnter nuumber of blocks per grid: ");
        scanf("%d",&B);
        if (T * B != N) printf("Error T x B != N, try again");
    } while (T * B != N);

    cudaEvent_t start, stop;    // using cuda events to measure time
    float elapsed_time_ms;    // which is applicable for asynchronous code also

    cudaMalloc((void**)&dev_a,N * sizeof(int));
    cudaMalloc((void**)&dev_b,N * sizeof(int));
```

```

    cudaMalloc((void**)&dev_c,N * sizeof(int));

    for(int i=0;i<N;i++) { // load arrays with some numbers
        a[i] = i;
        b[i] = i*1;
    }

    cudaMemcpy(dev_a, a , N*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b , N*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(dev_c, c , N*sizeof(int),cudaMemcpyHostToDevice);

    cudaEventCreate( &start ); // instrument code to measure start time
    cudaEventCreate( &stop );
    cudaEventRecord( start, 0 );

    add<<<B,T>>>(dev_a,dev_b,dev_c);

    cudaMemcpy(c,dev_c,N*sizeof(int),cudaMemcpyDeviceToHost);

    cudaEventRecord( stop, 0 ); // instrument code to measue end time
    cudaEventSynchronize( stop );
    cudaEventElapsedTime( &elapsed_time_ms, start, stop );

    for(int i=0;i<N;i++) {
        printf("%d+%d=%d\n",a[i],b[i],c[i]);
    }


    printf("Time to calculate results: %f ms.\n", elapsed_time_ms); // print out execution
time

    // clean up
    cudaFree(dev_a);
    cudaFree(dev_b);
    cudaFree(dev_c);

    cudaEventDestroy(start);
    cudaEventDestroy(stop);

    return 0;
}

```



Write a CUDA Program for Parallel Matrix Multiplication.

Program:

```
// Program for Matrix Addition in CUDA
```

```
// For Hadoop-CUDA Lab
```

```
#include <stdio.h>
```

```
#include <cuda.h>
```

```
#include <stdlib.h>
```

```
_global_ void gpu_matrixadd(int *a,int *b, int *c, int N) {
```

```
int col = threadIdx.x + blockDim.x * blockIdx.x;
int row = threadIdx.y + blockDim.y * blockIdx.y;
```

```
int index = row * N + col;
```

```
if(col < N && row < N)
c[index] = a[index]+b[index];
```

}

```
void cpu_matrixadd(int *a,int *b, int *c, int N) {
```

```
int index;
for(int col=0; col < N; col++)
    for(int row=0; row < N; row++) {
        index = row * N + col;
        c[index] = a[index]+b[index];
    }
```

}

```
int main(int argc, char *argv[]) {
```

```
char key;
```

```
int i, j;           // loop counters
```

```
int Grid_Dim_x=1, Grid_Dim_y=1; //Grid structure values
```

```
int Block_Dim_x=1, Block_Dim_y=1;           //Block structure values
```

```
int noThreads_x, noThreads_y;    // number of threads available in device, each
dimension
```

```
int noThreads_block;           // number of threads in a block
```

```
int N = 10;           // size of array in each dimension
```

```
int *a,*b,*c,*d;
```

```
int *dev_a, *dev_b, *dev_c;
```

```
int size; // number of bytes in arrays
```



```

        cudaEvent_t start, stop;                // using cuda events to measure time
        float elapsed_time_ms;                 // which is applicable for asynchronous code
also

/* -----ENTER INPUT PARAMETERS AND DATA -----*/

do { // loop to repeat complete program

    printf ("Device characteristics -- some limitations (compute capability 1.0)\n");
    printf ("        Maximum number of threads per block = 512\n");
    printf ("        Maximum sizes of x- and y- dimension of thread block = 512\n");
    printf ("        Maximum size of each dimension of grid of thread blocks = 65535\n");

    printf("Enter size of array in one dimension (square array), currently %d\n",N);
    scanf("%d",&N);

    do {

        printf("\nEnter number of blocks per grid in x dimension), currently %d :
",Grid_Dim_x);
        scanf("%d",&Grid_Dim_x);

        printf("\nEnter number of blocks per grid in y dimension), currently %d :
",Grid_Dim_y);
        scanf("%d",&Grid_Dim_y);

        printf("\nEnter number of threads per block in x dimension), currently %d :
",Block_Dim_x);
        scanf("%d",&Block_Dim_x);

        printf("\nEnter number of threads per block in y dimension), currently %d :
",Block_Dim_y);
        scanf("%d",&Block_Dim_y);

        noThreads_x = Grid_Dim_x * Block_Dim_x;                // number of threads in
x dimension
        noThreads_y = Grid_Dim_y * Block_Dim_y;                // number of threads in
y dimension

        noThreads_block = Block_Dim_x * Block_Dim_y;           // number of threads in
a block

```

```

        if (noThreads_x < N) printf("Error -- number of threads in x dimension less than
number of elements in arrays, try again\n");
        else if (noThreads_y < N) printf("Error -- number of threads in y dimension less
than number of elements in arrays, try again\n");
        else if (noThreads_block > 512) printf("Error -- too many threads in block, try
again\n");
        else printf("Number of threads not used = %d\n", noThreads_x * noThreads_y -
N * N);

    } while (noThreads_x < N || noThreads_y < N || noThreads_block > 512);

    dim3 Grid(Grid_Dim_x, Grid_Dim_x);          //Grid structure
    dim3 Block(Block_Dim_x,Block_Dim_y); //Block structure, threads/block limited by
specific device

    size = N * N * sizeof(int);                // number of bytes in total in arrays

    a = (int*) malloc(size);                    //this time use dynamically allocated memory for arrays
on host
    b = (int*) malloc(size);
    c = (int*) malloc(size);                    // results from GPU
    d = (int*) malloc(size);                    // results from CPU

    for(i=0;i < N;i++)                          // load arrays with some numbers
    for(j=0;j < N;j++) {
        a[i * N + j] = i;
        b[i * N + j] = i;
    }

/* ----- COMPUTATION DONE ON GPU ----- */

    cudaMalloc((void**)&dev_a, size);           // allocate memory on device
    cudaMalloc((void**)&dev_b, size);
    cudaMalloc((void**)&dev_c, size);

    cudaMemcpy(dev_a, a , size ,cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b , size ,cudaMemcpyHostToDevice);
    cudaMemcpy(dev_c, c , size ,cudaMemcpyHostToDevice);

    cudaEventCreate(&start);                     // instrument code to measure start time
    cudaEventCreate(&stop);

```

```

        cudaEventRecord(start, 0);
//      cudaEventSynchronize(start);    // Needed?

        gpu_matrixadd<<<Grid,Block>>>(dev_a,dev_b,dev_c,N);

        cudaMemcpy(c,dev_c, size ,cudaMemcpyDeviceToHost);

        cudaEventRecord(stop, 0);    // instrument code to measue end time
        cudaEventSynchronize(stop);
        cudaEventElapsedTime(&elapsed_time_ms, start, stop );

//      for(i=0;i < N;i++)
//      for(j=0;j < N;j++)
//      printf("%d+%d=%d\n",a[i * N + j],b[i * N + j],c[i * N + j]);

        printf("Time to calculate results on GPU: %f ms.\n", elapsed_time_ms); // print out
        execution time

/* ----- COMPUTATION DONE ON HOST CPU ----- */

        cudaEventRecord(start, 0);    // use same timing
//      cudaEventSynchronize(start);    // Needed?

        cpu_matrixadd(a,b,d,N);    // do calculation on host

        cudaEventRecord(stop, 0);    // instrument code to measue end time
        cudaEventSynchronize(stop);
        cudaEventElapsedTime(&elapsed_time_ms, start, stop );

        printf("Time to calculate results on CPU: %f ms.\n", elapsed_time_ms); // print out
        execution time

/* ----- check device creates correct results ----- */

        for(i=0;i < N*N;i++) {
                if (c[i] != d[i]) printf("**** ERROR in results, CPU and GPU create different
answers ***\n");
                break;
        }

        printf("\nEnter c to repeat, return to terminate\n");
        scanf("%c",&key);

```

```

scanf("%c",&key);

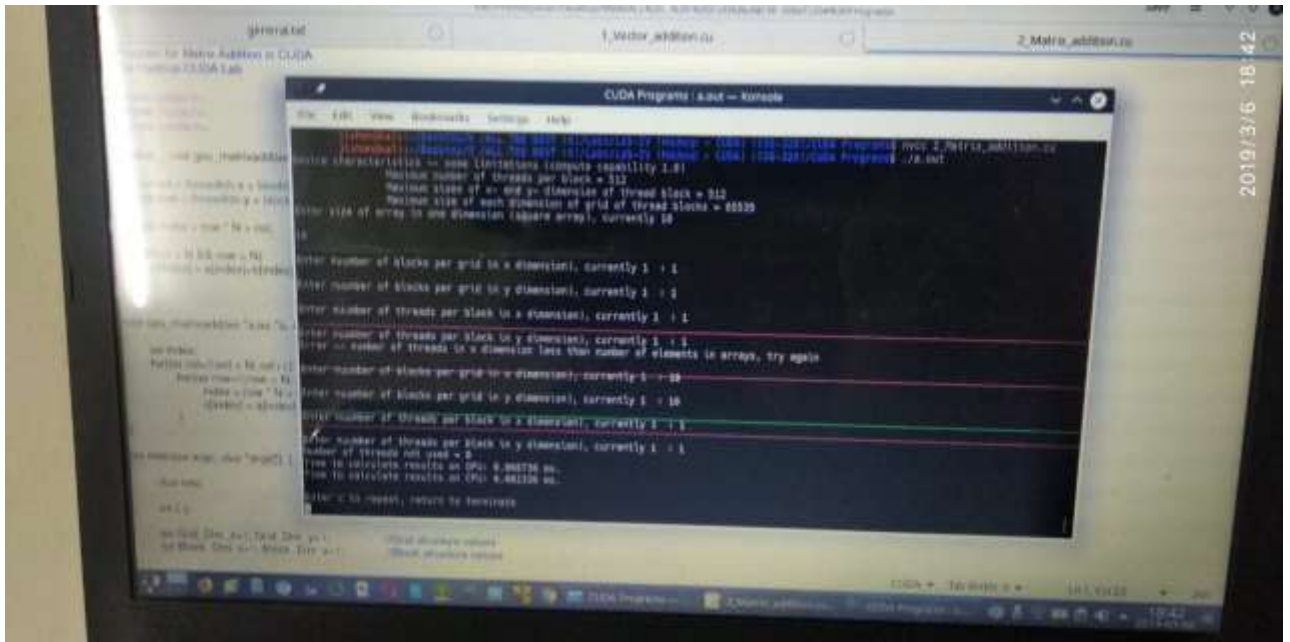
} while (key == 'c'); // loop of complete program

/* ----- clean up ----- */
free(a);
free(b);
free(c);
cudaFree(dev_a);
cudaFree(dev_b);
cudaFree(dev_c);

cudaEventDestroy(start);
cudaEventDestroy(stop);

return 0;
}

```



Write a CUDA Program for Parallel Binary Search.

Program: // Program for Parallel Binary Search in CUDA

// For Hadoop-CUDA Lab

```

#include <stdio.h>
#include <stdlib.h>
#include <cuda.h>
#include <cuda_runtime.h>

```

```
#include <assert.h>
```

```
_device_ int get_index_to_check(int thread, int num_threads, int set_size, int offset) {
```

```
    // Integer division trick to round up
```

```
    return (((set_size + num_threads) / num_threads) * thread) + offset;
```

```
}
```

```
_global_ void p_ary_search(int search, int array_length, int *arr, int *ret_val) {
```

```
    const int num_threads = blockDim.x * gridDim.x;
```

```
    const int thread = blockIdx.x * blockDim.x + threadIdx.x;
```

```
    //ret_val[0] = -1;
```

```
    //ret_val[1] = offset;
```

```
    int set_size = array_length;
```

```
    while(set_size != 0){
```

```
        // Get the offset of the array, initially set to 0
```

```
        int offset = ret_val[1];
```

```
        // I think this is necessary in case a thread gets ahead, and resets offset before
```

```
it's read
```

```
        // This isn't necessary for the unit tests to pass, but I still like it here
```

```
        __syncthreads();
```

```
        // Get the next index to check
```

```
        int index_to_check = get_index_to_check(thread, num_threads, set_size,  
offset);
```

```
        // If the index is outside the bounds of the array then lets not check it
```

```
        if (index_to_check < array_length){
```

```
            // If the next index is outside the bounds of the array, then set it to
```

```
maximum array size
```

```
            int next_index_to_check = get_index_to_check(thread + 1,  
num_threads, set_size, offset);
```

```

        if (next_index_to_check >= array_length){
            next_index_to_check = array_length - 1;
        }

        // If we're at the mid section of the array reset the offset to this index
        if (search > arr[index_to_check] && (search <
arr[next_index_to_check])) {
            ret_val[1] = index_to_check;
        }
        else if (search == arr[index_to_check]) {
            // Set the return var if we hit it
            ret_val[0] = index_to_check;
        }
    }

    // Since this is a p-ary search divide by our total threads to get the next set size
    set_size = set_size / num_threads;

    // Sync up so no threads jump ahead and get a bad offset
    __syncthreads();
}
}

```

```

int chop_position(int search, int *search_array, int array_length)
{
    // Get the size of the array for future use
    int array_size = array_length * sizeof(int);

    // Don't bother with small arrays
    if (array_size == 0) return -1;

    // Setup array to use on device
    int *dev_arr;
    cudaMalloc((void**)&dev_arr, array_size);

    // Copy search array values
    cudaMemcpy(dev_arr, search_array, array_size, cudaMemcpyHostToDevice);

    // return values here and on device
    int ret_val = (int)malloc(sizeof(int) * 2);
    ret_val[0] = -1; // return value
}

```

```

    ret_val[1] = 0; // offset
    array_length = array_length % 2 == 0 ? array_length : array_length - 1; // array size

    int          *dev_ret_val;
    cudaMalloc((void**)&dev_ret_val, sizeof(int) * 2);

    // Send in some intialized values
    cudaMemcpy(dev_ret_val, ret_val, sizeof(int) * 2, cudaMemcpyHostToDevice);

    // Launch kernel
    // This seems to be the best combo for p-ary search
    // Optimized around 10-15 registers per thread
    p_ary_search<<<16, 64>>>(search, array_length, dev_arr, dev_ret_val);

    // Get results
    cudaMemcpy(ret_val, dev_ret_val, 2 * sizeof(int), cudaMemcpyDeviceToHost);

    int ret = ret_val[0];

    printf("Ret Val %i   Offset %i\n", ret, ret_val[1]);

    // Free memory on device
    cudaFree(dev_arr);
    cudaFree(dev_ret_val);

    free(ret_val);

    return ret;
}

// Test region
static int * build_array(int length) {

    int ret_val = (int)malloc(length * sizeof(int));

    for (int i = 0; i < length; i++)
    {
        ret_val[i] = i * 2 - 1;
    }

    return ret_val;
}

```

```

static void test_array(int length, int search, int index) {

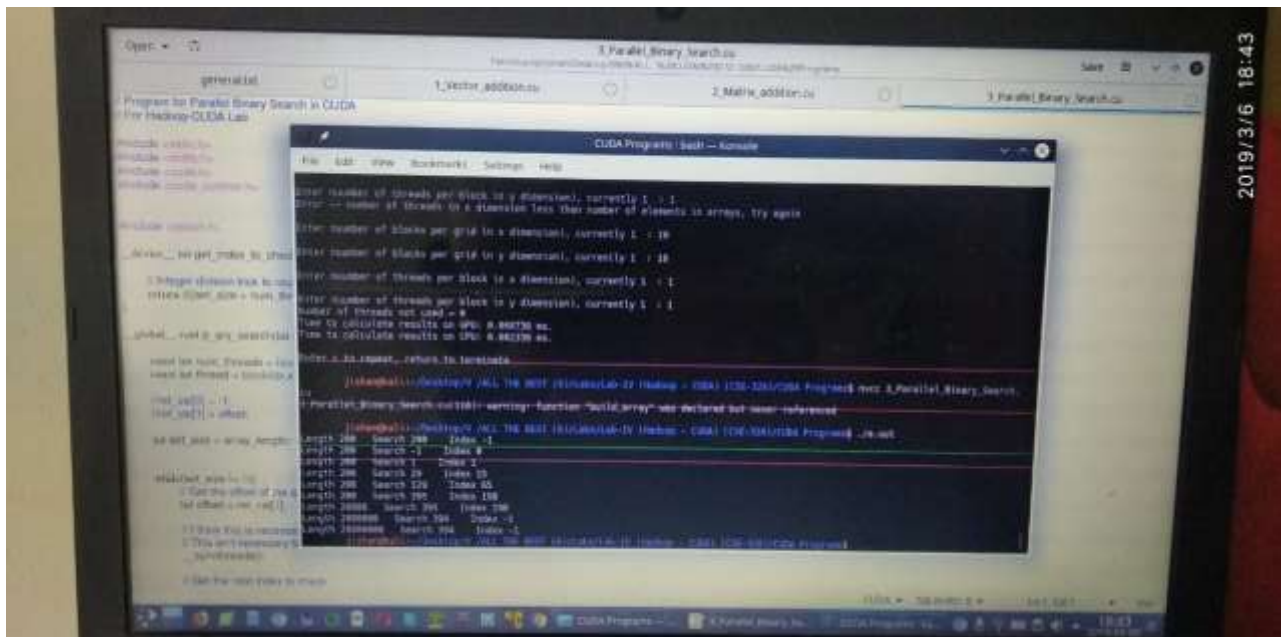
    printf("Length %i Search %i Index %i\n", length, search, index);
    // assert(index == chop_position(search, build_array(length), length) &&
    "test_small_array()");

}

static void test_arrays() {
    test_array(200, 200, -1);
    test_array(200, -1, 0);
    test_array(200, 1, 1);
    test_array(200, 29, 15);
    test_array(200, 129, 65);
    test_array(200, 395, 198);
    test_array(20000, 395, 198);
    test_array(2000000, 394, -1);
    test_array(20000000, 394, -1);
}

int main(){
    test_arrays();
}

```



HADOOP:

Write a Map Reduce program to calculate word frequency in given document

```
Program: import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
public class WordFrequency {
    public static void main(String [] args) throws Exception

    {

        Configuration conf=new Configuration();
        Job job=new Job(conf,"wordcount");
        job.setJarByClass(WordFrequency.class);
        job.setMapperClass(MapClass.class);
        job.setReducerClass(ReduceClass.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true)?0:1);

    }

    public static class MapClass extends Mapper<LongWritable, Text, Text, IntWritable>{

        public void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException
        {
            String line = value.toString();
            String[] words=line.split(" ");
```

```

        for(String word: words )
        {
            context.write(new Text(word.trim()), new IntWritable(1));
        }
    }
}

public static class ReduceClass extends Reducer<Text, IntWritable, Text, IntWritable>
{
    public void reduce(Text word, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException
    {
        int sum = 0;
        for(IntWritable value : values)
        {
            sum += value.get();
        }

        context.write(word, new IntWritable(sum));
    }
}

```

Calculate average marks of every student in the given dataset

```

Program: import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

```

public class AverageMarks {
    public static void main(String [] args) throws Exception

    {

        Configuration conf=new Configuration();

```

```

Job job=new Job(conf,"Average marks ");
job.setJarByClass(AverageMarks.class);
job.setMapperClass(MapClass.class);
job.setReducerClass(ReduceClass.class);
job.setOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);
job.setOutputValueClass(FloatWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true)?0:1);
}

```

```

public static class MapClass extends Mapper<LongWritable, Text, Text, IntWritable>{

    public void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException
    {
        String line = value.toString();
        String[] Words=line.split(",");
        if(Words.length>1)
        {
            context.write(new Text(Words[0]), new
IntWritable(Integer.parseInt(Words[1])));
        }
    }
}

public static class ReduceClass extends Reducer<Text, IntWritable, Text, FloatWritable>
{
    public void reduce(Text word, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException
    {
        int total= 0;
        int count=0;
        for(IntWritable value : values)
        {
            total += value.get();
            count+=1;
        }

        context.write(word, new FloatWritable(total/count));
    }
}

```

```

    }
}

```

Write a program to illustrate use of partitioner in Map Reduce

```

Program: import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.Partitioner;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class Driver{
    public static void main(String [] args) throws Exception{

        Configuration conf =new Configuration();
        Job job=new Job(conf,"Implement Partitioner");
        job.setJarByClass(Driver.class);
        job.setMapperClass(MapClass.class);
        job.setPartitionerClass(PartitionClass.class);
        job.setReducerClass(ReduceClass.class);
        job.setNumReduceTasks(6);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(Text.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true)?0:1);
    }

    public static class MapClass extends Mapper<LongWritable, Text, Text, Text>{

        public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException
        {
            String line = value.toString();

```

```

        String[] Words=line.split(",");
        if(Words.length>1)
        {
            context.write(new Text(Words[0]), new Text(Words[1]));
        }
    }
}

public static class PartitionClass extends Partitioner<Text,Text>{
    public int getPartition(Text key,Text value,int numReduceTask)
    {
        if(numReduceTask==0)
            return 0;
        if(key.equals(new Text("Cse1")))
            return 0;
        else if(key.equals(new Text("Mech1")))
            return 1%numReduceTask;
        else if(key.equals(new Text("Cse2")))
            return 2%numReduceTask;
        else if(key.equals(new Text("Mech2")))
            return 3%numReduceTask;
        else if (key.equals(new Text("Civil")))
            return 4%numReduceTask;
        else
            return 5%numReduceTask;
    }
}

public static class ReduceClass extends Reducer<Text,Text,Text,Text>{
    public void reduce(Text key,Iterable<Text> Values,Context context)throws
IOException, InterruptedException{
        int sum=0;
        for(Text val:Values)
        {
            context.write(key,val);
        }
    }
}
}

```